

Contents

1	Mason and Grijalva	2
2	Wang and Hong	2
3	Han et al.	4
4	Yu et al.	4
5	Sierla et al.	6
6	Fu et al.	6
7	Shaqour and Hagishima	8
8	Zhang et al.	8
9	Al-Ani and Das	10
10	Weinberg et al.	10
11	Chatterjee and Khovalyg	12
12	Al Sayed et al.	12
13	Pinthurat et al.	14
14	Latoń et al.	14
15	Huynh et al.	16
16	Kannari et al.	16
17	Michailidis et al.	18
18	Li et al.	18
19	Togashi	20
20	Lei et al.	20
21	Mazhar et al.	22
22	Mohsenpour and Xing	23

1 Mason and Grijalva

Table 1: MRQ5 recoded audit table for Mason and Grijalva (2019). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
MasonGrijalva2019	2	2	2	2	The review discusses validation maturity mainly through the dominance of simulated building energy-management problems, the dependence of RL evaluation on accurate simulator design, and the need for real-world-representative data. Deployment-readiness treatment receives score 2 because the paper explains that direct online learning in physical buildings would involve costly trial-and-error exploration, making simulation pretraining and later physical refinement important. Simulation-to-real discussion receives score 2 because the paper explicitly connects simulator fidelity, physical-building implementation, pretraining, and transfer learning. Gap-signal visibility receives score 2 because the paper clearly identifies simulation dependence and transfer to broader deployment as limitations, although it does not provide a full validation taxonomy or staged deployment ladder.

Table 2: MRQ5 validation-modality scores for Mason and Grijalva (2019), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Mason and Grijalva (2019)	3	0	0	1	1

Table 3: MRQ5 validation-rigor scores for Mason and Grijalva (2019), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Mason and Grijalva (2019)	1	2	2	1	1

Table 4: MRQ5 gap-signal scores for Mason and Grijalva (2019), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Mason and Grijalva (2019)	2	1	2	1	0

Table 5: MRQ5 summary table for Mason and Grijalva (2019), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Mason and Grijalva (2019)	Strongly simulation-dominant with early deployment caution	Validation is framed mainly as simulated training followed, where possible, by later physical implementation or refinement. The key concern is that direct real-building learning is too costly during early exploration, so credible simulation and transfer become important.	The survey recognizes the simulation-to-real problem clearly, but it does not yet offer a well-developed validation ladder across testbeds, pilots, field trials, and long-term deployment maturity stages.

Table 6: One-line MRQ5 answer for Mason and Grijalva (2019), after recoding.

survey_id	MRQ5 answer
MasonGrijalva2019	Mason and Grijalva (2019) treat validation maturity mainly as a shift from simulation-trained RL toward limited physical implementation, emphasizing the need for accurate simulators, warning against costly direct exploration in real buildings, and pointing to transfer learning as a future route to wider deployment.

2 Wang and Hong

Table 7: MRQ5 recoded audit table for Wang and Hong (2020). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
WangHong2020	3	3	3	3	The review structurally distinguishes simulation environments from real buildings and reports that real-building deployment remains rare. Validation-maturity treatment receives score 3 because simulation-heavy evidence and limited real-building implementation are central to the paper’s assessment of RL for building controls. Deployment-readiness treatment receives score 3 because the paper explicitly organizes deployment barriers around time- and data-intensive training, control security and robustness, and limited generalization. Simulation-to-real discussion receives score 3 because the paper links simulated evaluation, real-building deployment limits, transfer learning, and the need for shared test infrastructure. Gap-signal visibility receives score 3 because open-source testbeds, datasets, benchmarking, and transfer learning are presented as structural needs for improving comparability and deployment readiness.

Table 8: MRQ5 validation-modality scores for Wang and Hong (2020), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Wang and Hong (2020)	3	2	0	1	2

Table 9: MRQ5 validation-rigor scores for Wang and Hong (2020), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Wang and Hong (2020)	2	3	3	2	1

Table 10: MRQ5 gap-signal scores for Wang and Hong (2020), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Wang and Hong (2020)	2	2	2	1	1

Table 11: MRQ5 summary table for Wang and Hong (2020), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Wang and Hong (2020)	Simulation-dominant with explicit real-deployment barriers	Validation is framed through a strong simulation-versus-real-building contrast, with actual deployment treated as limited and blocked by training burden, security and robustness concerns, and weak generalization.	The survey gives a clear deployment-readiness critique rather than only a list of validation modes. Its main MRQ5 value is in diagnosing why simulated RL success has not translated easily into real-building adoption and why open testbeds, datasets, and transfer learning are needed.

Table 12: One-line MRQ5 answer for Wang and Hong (2020), after recoding.

survey_id	MRQ5 answer
WangHong2020	Wang and Hong (2020) define validation maturity mainly through the contrast between simulation-heavy RL research and rare real-building deployment, arguing that deployment readiness is constrained by expensive training, insufficient robustness and security, and weak generalization, with transfer learning and shared test infrastructure proposed as partial remedies.

3 Han et al.

Table 13: MRQ5 recoded audit table for Han et al. (2021). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
HanEtAl2021	3	2	3	3	The review treats validation maturity as a central difficulty for occupant-behavior-oriented RL in building control. Validation-maturity treatment receives score 3 because the paper explicitly states that there is no general scientific standard for validating models when multiple occupant behaviors are involved. Deployment-readiness treatment receives score 2 because the paper discusses implementation difficulty, training requirements, robustness, and real implementations as a way to validate and improve methods, but deployment is not the main organizing axis. Simulation-to-real discussion receives score 3 because context-dependent occupant behavior and transfer from one setting to a more complex scenario are central validation concerns. Gap-signal visibility receives score 3 because lack of validation standards, transfer difficulty, context dependence, and generalization limits are explicit limitations in the paper.

Table 14: MRQ5 validation-modality scores for Han et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Han et al. (2021)	3	1	0	1	2

Table 15: MRQ5 validation-rigor scores for Han et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Han et al. (2021)	3	2	3	1	1

Table 16: MRQ5 gap-signal scores for Han et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Han et al. (2021)	1	1	3	1	1

Table 17: MRQ5 summary table for Han et al. (2021), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Han et al. (2021)	Simulation-oriented with explicit concern about transfer and validation ambiguity	Validation is framed mainly through the absence of clear validation standards, the context dependence of occupant behavior, and the difficulty of transferring knowledge from one setting to another more complex one.	The survey contributes to MRQ5 less by building a staged validation ladder and more by showing that occupant-centered RL faces a poorly standardized and transfer-sensitive path from simulated studies toward credible real-world use.

Table 18: One-line MRQ5 answer for Han et al. (2021), after recoding.

survey_id	MRQ5 answer
HanEtAl2021	Han et al. (2021) treat validation maturity mainly as a problem of weak validation standards, context-dependent occupant behavior, and limited transfer from simulated or simplified settings to more complex real-world environments, while only lightly touching actual implementation as a way to validate and improve RL methods.

4 Yu et al.

Table 19: MRQ5 recoded audit table for Yu et al. (2021). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
YuEtAl2021	3	3	3	3	The review evaluates DRL methods in terms of convergence, effectiveness, scalability, and robustness, while also distinguishing model-free and model-based DRL in terms of interaction burden, exploration cost, and deployment feasibility. Validation-maturity treatment receives score 3 because implementation mode is structurally visible through the model-free/model-based comparison and tabulated practical-implementation evidence. Deployment-readiness treatment receives score 3 because the paper explicitly argues that fully exploring the state space in real buildings is impractical due to high cost and long training time, and it presents deployment-oriented pipelines involving building energy modeling, calibration, offline training, and real deployment. Simulation-to-real discussion receives score 3 because practical readiness is framed through calibrated models, offline training, domain randomization, transfer, and movement from virtual environments into real buildings. Gap-signal visibility receives score 3 because costly real-building exploration, simulation-heavy evaluation, and the need for model-based routes to deployment are central limitations and future-facing concerns in the survey.

Table 20: MRQ5 validation-modality scores for Yu et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Simulation Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Yu et al. (2021)	3	1	1	2

Table 21: MRQ5 validation-rigor scores for Yu et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Yu et al. (2021)	3	3	3	1	2

Table 22: MRQ5 gap-signal scores for Yu et al. (2021), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Yu et al. (2021)	2	1	2	1	1

Table 23: MRQ5 summary table for Yu et al. (2021), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Yu et al. (2021)	Simulation-dominant but deployment-structured	Validation is framed mainly through the contrast between model-free DRL, which is mostly simulation-evaluated, and model-based DRL, which is presented as the main route toward practical deployment through modeling, calibration, offline training, and transfer to real environments.	The survey answers MRQ5 less by offering a full validation ladder and more by identifying a practical bridge from simulation to deployment, where model-based DRL serves as the main pathway for reducing costly and risky real-building exploration.

Table 24: One-line MRQ5 answer for Yu et al. (2021), after recoding.

survey_id	MRQ5 answer
YuEtAl2021	Yu et al. (2021) treat validation maturity mainly through the difference between simulation-heavy model-free DRL and more deployment-oriented model-based DRL, arguing that real-building exploration is too costly and slow for direct training and that practical readiness therefore depends on calibrated models, offline training, and transfer into real environments.

5 Sierla et al.

Table 25: MRQ5 recoded audit table for Sierla et al. (2022). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
SierlaEtAl2022	3	3	2	3	The review treats validation maturity not only as a question of where an RL controller is tested, but also as a question of whether the control formulation is physically meaningful and deployable in HVAC systems. Validation-maturity treatment receives score 3 because the paper structurally analyzes action formulations, HVAC abstraction levels, and whether RL actions correspond to actuator commands, control-loop setpoints, or supervisory variables. Deployment-readiness treatment receives score 3 because the feasibility of interfacing RL outputs with actual HVAC equipment is central to the paper’s critique. Simulation-to-real discussion receives score 2 because the paper distinguishes physical HVAC systems from virtual environments such as building simulators and data-driven models, but it does not provide a full sim-to-real pipeline. Gap-signal visibility receives score 3 because vague action mappings, excessive abstraction, and unclear deployment interfaces are structural barriers identified by the review.

Table 26: MRQ5 validation-modality scores for Sierla et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Simulation Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Sierla et al. (2022)	2	1	0	1

Table 27: MRQ5 validation-rigor scores for Sierla et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Sierla et al. (2022)	3	3	2	1	1

Table 28: MRQ5 gap-signal scores for Sierla et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Sierla et al. (2022)	1	1	2	3	1

Table 29: MRQ5 summary table for Sierla et al. (2022), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Sierla et al. (2022)	Deployment-conscious and interface-oriented, but still largely virtual-environment aware	Validation is framed mainly through whether the RL action space and abstraction level are realistic enough to interface with actual HVAC systems, rather than through a formal ladder from simulation to field deployment.	The survey contributes to MRQ5 by showing that validation maturity depends not only on where the controller is tested, but also on whether the control formulation itself is physically meaningful and deployable in a BACS/HVAC context.

Table 30: One-line MRQ5 answer for Sierla et al. (2022), after recoding.

survey__id	MRQ5 answer
SierlaEtAl2022	Sierla et al. (2022) treat validation maturity mainly as a question of deployable control formulation, arguing that RL studies should be judged by whether their action spaces and abstraction levels can interface credibly with real HVAC systems, not just by simulated optimization performance.

6 Fu et al.

“`latex`

Table 31: MRQ5 recoded audit table for Fu et al. (2022). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
FuEtAl2022	2	3	3	3	The review discusses validation maturity mainly through the practical difficulty of applying RL in real buildings and the continuing reliance on simulation platforms. Validation-maturity treatment receives score 2 because the paper discusses simulation platforms and some practical experiments or real-world implementations, but does not build a full validation taxonomy. Deployment-readiness treatment receives score 3 because the paper explicitly identifies real-building training as time-consuming, costly, inconvenient for occupants, potentially damaging to equipment, associated with uncontrollable costs, and affected by security concerns. Simulation-to-real discussion receives score 3 because the difference between simulated and real environments is named as a major obstacle. Gap-signal visibility receives score 3 because deployment cost, safety/security risk, training burden, environment modeling, and simulation-real mismatch are explicit barriers in the review’s challenge discussion.

Table 32: MRQ5 validation-modality scores for Fu et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Simulation Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Fu et al. (2022)	3	1	1	0

Table 33: MRQ5 validation-rigor scores for Fu et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Fu et al. (2022)	1	3	3	3	1

Table 34: MRQ5 gap-signal scores for Fu et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Fu et al. (2022)	2	1	3	2	1

Table 35: MRQ5 summary table for Fu et al. (2022), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Fu et al. (2022)	Simulation-dominant with explicit real-deployment obstacles	Validation is framed mainly through the practical difficulty of moving RL from simulation into real buildings, especially because trial-and-error training in real environments is slow, costly, risky for occupants and equipment, and affected by security concerns.	The survey contributes to MRQ5 mainly by showing that deployment readiness is constrained less by lack of algorithmic promise than by the operational burden of real-world training and by unresolved differences between simulated and real environments.

Table 36: One-line MRQ5 answer for Fu et al. (2022), after recoding.

survey_id	MRQ5 answer
FuEtAl2022	Fu et al. (2022) treat validation maturity mainly as a practical deployment problem, arguing that RL remains heavily simulation-based because direct training in real buildings is time-consuming, costly, potentially unsafe for occupants and equipment, and affected by a persistent gap between simulated and real environments.

7 Shaqour and Hagishima

Table 37: MRQ5 recoded audit table for Shaqour and Hagishima (2022). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
ShaqourHagishima2022	3	3	3	3	The review treats validation maturity as a structurally visible issue in DRL-based BEMS research. Validation-maturity treatment receives score 3 because the paper reports the scarcity of real implementations and validations while also organizing reviewed studies by building type and validation setting. Deployment-readiness treatment receives score 3 because the paper emphasizes that pre-training or offline training before real-world deployment is critical, since untrained DRL agents are not suitable for direct deployment. Simulation-to-real discussion receives score 3 because the paper connects data hunger, offline training, scarce real-world validation, and the need to move beyond purely simulated evaluation. Gap-signal visibility receives score 3 because limited real-system validation, data requirements, and deployment-readiness barriers are explicit recurring limitations, even though laboratory, university-building, school, and other real-system examples are beginning to appear.

Table 38: MRQ5 validation-modality scores for Shaqour and Hagishima (2022), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Shaqour and Hagishima (2022)	3	2	2	1	1

Table 39: MRQ5 validation-rigor scores for Shaqour and Hagishima (2022), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Shaqour and Hagishima (2022)	3	3	3	1	2

Table 40: MRQ5 gap-signal scores for Shaqour and Hagishima (2022), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Shaqour and Hagishima (2022)	2	2	2	1	2

Table 41: MRQ5 summary table for Shaqour and Hagishima (2022), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Shaqour and Hagishima (2022)	Simulation-dominant with explicit concern about real validation scarcity	Validation is framed mainly through the need for sufficient offline training before deployment and through the observation that real implementations remain rare, although laboratory, school, university-building, and other real-system validations provide early non-simulation evidence.	The survey contributes to MRQ5 by showing that deployment readiness in DRL-based BEMS is strongly constrained by data availability and pre-training requirements, while also indicating that intermediate real-system validation forms are beginning to appear.

Table 42: One-line MRQ5 answer for Shaqour and Hagishima (2022), after recoding.

survey_id	MRQ5 answer
ShaqourHagishima2022	Shaqour and Hagishima (2022) treat validation maturity mainly as a problem of moving data-hungry DRL methods beyond simulation through sufficient offline training and scarce real-system validation, while noting that laboratory, school, university-building, and other real-system experiments provide early but still limited evidence of deployment readiness.

8 Zhang et al.

Table 43: MRQ5 recoded audit table for Zhang et al. (2022). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
ZhangEtAl2022	2	3	3	3	The review discusses validation maturity mainly through real-world feasibility rather than through a formal validation taxonomy. Validation-maturity treatment receives score 2 because the paper discusses the limitations of direct RL exploration in buildings, but does not structure the review around validation stages. Deployment-readiness treatment receives score 3 because real-world applicability is explicitly framed around data efficiency, safety, robust adaptability, high economic cost from bad exploratory actions, and long real-world training time. Simulation-to-real discussion receives score 3 because the paper presents model-based RL, learning from historical data, and MPC/RL combinations as practical routes for reducing unsafe or inefficient real-building exploration. Gap-signal visibility receives score 3 because safety, data efficiency, robust adaptability, and real-world feasibility are explicit deployment barriers and future-facing limitations in the review.

Table 44: MRQ5 validation-modality scores for Zhang et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Zhang et al. (2022)	2	1	0	1	2

Table 45: MRQ5 validation-rigor scores for Zhang et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Zhang et al. (2022)	1	3	3	3	1

Table 46: MRQ5 gap-signal scores for Zhang et al. (2022), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Zhang et al. (2022)	1	1	2	2	0

Table 47: MRQ5 summary table for Zhang et al. (2022), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Zhang et al. (2022)	Deployment-conscious and hybrid-method oriented	Validation is framed mainly through real-world feasibility, with direct RL deployment treated as costly and risky, and model-based or MPC-assisted approaches presented as more practical routes toward usable building control.	The survey contributes to MRQ5 by recasting the simulation-to-real gap as a practical control-design problem centered on data efficiency, safety, and robust adaptability rather than as a formal maturity ladder across validation stages.

Table 48: One-line MRQ5 answer for Zhang et al. (2022), after recoding.

survey_id	MRQ5 answer
ZhangEtAl2022	Zhang et al. (2022) treat validation maturity mainly as a question of real-world feasibility, arguing that direct RL in buildings is limited by expensive exploration, safety concerns, and weak robust adaptability, so stronger deployment readiness depends on model-based RL, historical data use, and MPC-assisted or hybrid control designs.

Table 49: MRQ5 recoded audit table for Al-Ani and Das (2022). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
AlAniDas2022	1	1	1	1	The paper is a broad review of reinforcement learning theory and applications in home energy management systems. Validation-maturity treatment receives score 1 because the paper mentions that RL research in HEMS is still at an early stage and that actual HEMS deployment remains limited, but validation maturity does not structure the review. Deployment-readiness treatment receives score 1 because the paper distinguishes offline model-based training from online interaction with the real physical environment, but does not develop a detailed deployment-readiness framework. Simulation-to-real discussion receives score 1 because the paper notes the use of HEMS models and simulation-based training, but does not provide a sustained simulation-to-real analysis. Gap-signal visibility receives score 1 because limited deployment and early-stage maturity are acknowledged, while the paper’s main structure remains focused on RL methods, HEMS application classes, and performance metrics rather than validation or deployment gaps.

Table 50: MRQ5 validation-modality scores for Al-Ani and Das (2022), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Al-Ani and Das (2022)	2	0	0	1	0

Table 51: MRQ5 validation-rigor scores for Al-Ani and Das (2022), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Al-Ani and Das (2022)	1	1	1	0	1

Table 52: MRQ5 gap-signal scores for Al-Ani and Das (2022), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Al-Ani and Das (2022)	1	1	0	0	0

Table 53: MRQ5 summary table for Al-Ani and Das (2022), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Al-Ani and Das (2022)	Broad and lightly deployment-aware, but still model/simulation-oriented	Validation is framed mainly at a high level through the distinction between offline model-based training and online interaction with the real environment, together with the observation that actual HEMS deployment remains limited in prior literature.	The survey contributes to MRQ5 mainly by confirming that RL for HEMS remains early-stage and still largely trained through models, rather than by developing a strong validation-maturity ladder or a systematic simulation-to-real critique.

Table 54: One-line MRQ5 answer for Al-Ani and Das (2022), after recoding.

survey_id	MRQ5 answer
AlAniDas2022	Al-Ani and Das (2022) treat validation maturity mainly as an early-stage HEMS literature issue, acknowledging limited reported deployment and widespread model-based training, but focusing more on broad method coverage, HEMS application classes, and performance metrics than on a detailed deployment-readiness or simulation-to-real framework.

Table 55: MRQ5 recoded audit table for Weinberg et al. (2023). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
WeinbergEtAl2023	2	3	3	3	The review is a computer-science-oriented survey of reinforcement learning for building energy systems. Validation-maturity treatment receives score 2 because the paper does not build a staged validation taxonomy, but discusses practical readiness through sample efficiency, environment-interaction burden, transfer, scalability, and risky exploration. Deployment-readiness treatment receives score 3 because deployment readiness is structurally tied to the paper’s computer-science challenge framework, including sample efficiency, transfer between buildings, scalability, multi-agent control, communication, and system design. Simulation-to-real discussion receives score 3 because transferability, sample efficiency, environment-interaction burden, and model-based reinforcement learning are central to how the paper frames practical use. Gap-signal visibility receives score 3 because the paper is explicitly organized around research gaps and future directions in these computer-science challenges.

Table 56: MRQ5 validation-modality scores for Weinberg et al. (2023), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Weinberg et al. (2023)	1	0	0	0	2

Table 57: MRQ5 validation-rigor scores for Weinberg et al. (2023), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Weinberg et al. (2023)	1	3	3	2	0

Table 58: MRQ5 gap-signal scores for Weinberg et al. (2023), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Weinberg et al. (2023)	1	0	3	2	0

Table 59: MRQ5 summary table for Weinberg et al. (2023), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Weinberg et al. (2023)	Technical-readiness oriented, with strong emphasis on transfer and sample efficiency	Validation is framed mainly through the technical conditions needed for practical reinforcement learning in building energy systems, especially sample efficiency, reduced environment interaction, transfer of policies between buildings, scalability, and multi-agent coordination, rather than through an explicit ladder from simulation to field deployment.	The survey contributes to MRQ5 by showing that deployment maturity depends heavily on whether reinforcement learning can become data-efficient, transferable, scalable, and safe enough for real building energy system operation, even though the paper is less focused on direct field-validation categories.

Table 60: One-line MRQ5 answer for Weinberg et al. (2023), after recoding.

survey__id	MRQ5 answer
WeinbergEtAl2023	Weinberg et al. (2023) treat validation maturity mainly as a problem of technical readiness for practical building energy system control, arguing that wider real-world deployment requires stronger sample efficiency, reduced risky environment interaction, transferability across buildings, scalability, and multi-agent coordination, rather than merely more simulation studies.

11 Chatterjee and Khovalyg

Table 61: MRQ5 recoded audit table for Chatterjee and Khovalyg (2023). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
ChatterjeeKhovalyg2023	3	3	3	3	The review treats validation and deployment maturity as structurally visible issues in reinforcement-learning-based HVAC control. Validation-maturity treatment receives score 3 because the paper summarizes simulation software, building model scale, building type, and real-building implementation status. Deployment-readiness treatment receives score 3 because the conclusion explicitly states that reinforcement-learning controllers remain largely in the research-and-development stage, with limited acceptance in actual buildings. Simulation-to-real discussion receives score 3 because the paper contrasts simulation-trained controllers with limited real-building implementation and identifies training burden, control security, and transfer learning as major deployment obstacles. Gap-signal visibility receives score 3 because the paper explicitly calls for data-rich, open-source, interoperable virtual testbeds for cross-study validation and benchmarking.

Table 62: MRQ5 validation-modality scores for Chatterjee and Khovalyg (2023), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Chatterjee and Khovalyg (2023)	3	1	1	2	1

Table 63: MRQ5 validation-rigor scores for Chatterjee and Khovalyg (2023), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Chatterjee and Khovalyg (2023)	3	3	3	2	2

Table 64: MRQ5 gap-signal scores for Chatterjee and Khovalyg (2023), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Chatterjee and Khovalyg (2023)	2	2	3	2	2

Table 65: MRQ5 summary table for Chatterjee and Khovalyg (2023), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Chatterjee and Khovalyg (2023)	Simulation-dominant but explicitly deployment-aware	Validation is framed through a strong contrast between simulation-trained reinforcement-learning controllers and their still-limited use in actual buildings, with deployment readiness tied to training burden, control security, transfer learning, and better shared virtual test infrastructure.	The survey contributes to MRQ5 by making the reality-gap problem unusually explicit: most reinforcement-learning studies are still simulation-trained, only a minority reach real buildings, and stronger deployment maturity will require safer, more data-efficient, and more transferable control pipelines.

Table 66: One-line MRQ5 answer for Chatterjee and Khovalyg (2023), after recoding.

survey_id	MRQ5 answer
ChatterjeeKhovalyg2023	Chatterjee and Khovalyg (2023) treat validation maturity mainly as the problem of moving reinforcement-learning controllers beyond simulation into real buildings, arguing that the field remains largely simulation-trained and that wider deployment is still constrained by training cost, control security, limited transferability, and the lack of shared validation infrastructure.

12 Al Sayed et al.

Table 67: MRQ5 recoded audit table for Al Sayed et al. (2024). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
AlSayedEtAl2024	3	3	3	3	The review treats virtual-to-real deployment as a structuring concern for reinforcement-learning-based HVAC control. Validation-maturity treatment receives score 3 because the paper reports limited real-building deployment and structurally reviews simulation environments, model environments, and training episode design. Deployment-readiness treatment receives score 3 because the paper identifies unsafe direct real use, insufficient exogenous-state diversification, and costly retraining or fine-tuning as central barriers to practical deployment. Simulation-to-real discussion receives score 3 because the paper directly addresses virtual training environments, real-world distribution shift, changing occupancy schedules, electricity prices, and weather, and the risk that learned virtual distributions become outdated before deployment. Gap-signal visibility receives score 3 because weak virtual-to-real transfer, limited training diversity, retraining cost, and the need for meta-reinforcement-learning adaptation are structural limitations and future directions in the review.

Table 68: MRQ5 validation-modality scores for Al Sayed et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Simulation Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Al Sayed et al. (2024)	3	0	1	1

Table 69: MRQ5 validation-rigor scores for Al Sayed et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Al Sayed et al. (2024)	3	3	3	3	1

Table 70: MRQ5 gap-signal scores for Al Sayed et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Al Sayed et al. (2024)	2	2	3	3	2

Table 71: MRQ5 summary table for Al Sayed et al. (2024), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Al Sayed et al. (2024)	Simulation-dominant but strongly deployment-aware	Validation is framed mainly through the gap between virtual training environments and non-stationary real-world deployment, with readiness depending on whether training episodes capture sufficient diversity in exogenous state components and whether retraining can be reduced.	The survey contributes to MRQ5 by showing that deployment maturity is blocked less by lack of algorithmic promise than by unsafe direct transfer from simulation, weak training diversity, distribution shift, and expensive retraining under real-world change.

Table 72: One-line MRQ5 answer for Al Sayed et al. (2024), after recoding.

survey__id	MRQ5 answer
AlSayedEtAl2024	Al Sayed et al. (2024) treat validation maturity mainly as a problem of safely moving simulation-trained HVAC controllers into non-stationary real buildings, arguing that limited real deployment, insufficiently diverse training episodes, distribution shift, and costly retraining make generalization the main barrier, with meta-reinforcement learning proposed as a promising remedy.

Table 73: MRQ5 recoded audit table for Pinthurat et al. (2024). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
PinthuratEtAl2024	3	3	3	3	The review treats real-world testing and validation as a structuring future-research need for reinforcement-learning-based smart home energy management systems with energy storage. Validation-maturity treatment receives score 3 because the paper includes a dedicated discussion of real-world testing and validation as necessary for assessing performance, adaptability, scalability, and practical viability. Deployment-readiness treatment receives score 3 because the paper explicitly connects deployment readiness to field trials under real user behavior, weather, and grid-dynamic conditions. Simulation-to-real discussion receives score 3 because validation in real environments is presented as necessary to refine algorithms, validate models, optimize parameters, and bridge theory and practice. Gap-signal visibility receives score 3 because lack of real-world testing, field validation, and practical deployment evidence is an explicit future-research direction, alongside concerns about data intensity, cost, safety, privacy, and security.

Table 74: MRQ5 validation-modality scores for Pinthurat et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Pinthurat et al. (2024)	2	0	1	2	2

Table 75: MRQ5 validation-rigor scores for Pinthurat et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Pinthurat et al. (2024)	1	3	3	2	1

Table 76: MRQ5 gap-signal scores for Pinthurat et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Pinthurat et al. (2024)	1	3	3	2	2

Table 77: MRQ5 summary table for Pinthurat et al. (2024), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Pinthurat et al. (2024)	Future-oriented and deployment-aware, but still field-evidence limited	Validation is framed mainly as a missing practical step: the paper emphasizes the need for real-world testing and field validation to examine adaptability, scalability, user behavior, weather variation, grid dynamics, and practical viability.	The survey contributes to MRQ5 by arguing that reinforcement-learning-based smart-home energy management with storage still needs much stronger real-world validation before deployment readiness can be claimed with confidence. Its strongest contribution is not a validation taxonomy, but a clear future agenda for field testing and theory-to-practice validation.

Table 78: One-line MRQ5 answer for Pinthurat et al. (2024), after recoding.

survey__id	MRQ5 answer
PinthuratEtAl2024	Pinthurat et al. (2024) treat validation maturity mainly as a gap between promising RL-based smart-home energy management methods and the still-limited real-world testing needed to verify performance, adaptability, scalability, and practical viability under real user, weather, and grid conditions.

Table 79: MRQ5 recoded audit table for Latoń et al. (2024). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
LatonEtAl2024	3	3	3	3	The review treats validation maturity and practical implementation as structurally visible issues for reinforcement learning and deep reinforcement learning in home energy management systems. Validation-maturity treatment receives score 3 because the paper explicitly contrasts theoretical or simulation-based research with the need for practical case studies and real-world implementations, and its tabulated review shows simulation-heavy verification, limited real implementation, and use of real-world data in only some cases. Deployment-readiness treatment receives score 3 because the paper discusses practical implementation barriers including data intensity, safety and reliability, scalability, heterogeneous infrastructure, privacy, and security. Simulation-to-real discussion receives score 3 because the paper frames maturity as a theory-to-practice gap that requires testing under real conditions and translating simulation results into operational evidence. Gap-signal visibility receives score 3 because lack of real applications, lack of detailed case studies, limited tangible operational impact, security/privacy issues, and scalability constraints are explicit gaps and future needs.

Table 80: MRQ5 validation-modality scores for Latoń et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Latoń et al. (2024)	3	0	1	1	2

Table 81: MRQ5 validation-rigor scores for Latoń et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Latoń et al. (2024)	3	3	3	3	1

Table 82: MRQ5 gap-signal scores for Latoń et al. (2024), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Latoń et al. (2024)	3	3	2	2	3

Table 83: MRQ5 summary table for Latoń et al. (2024), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Latoń et al. (2024)	Simulation-dominant but explicitly practice-oriented	Validation is framed mainly through the gap between simulation-heavy RL/DRL research and the need for practical case studies and real-world implementations, with deployment readiness linked to data pipelines, safety, reliability, scalability, privacy, security, and heterogeneous system integration.	The survey contributes to MRQ5 by showing that maturity is limited less by lack of promising algorithms than by the shortage of real-world case evidence and by the broader practical conditions required for secure, reliable, scalable, and operational deployment.

Table 84: One-line MRQ5 answer for Latoń et al. (2024), after recoding.

survey_id	MRQ5 answer
LatonEtAl2024	Latoń et al. (2024) treat validation maturity mainly as a theory-to-practice gap, arguing that RL/DRL for HEMS remains largely simulation-based and that stronger real-world case studies, together with scalable, safe, reliable, and privacy-preserving deployment infrastructures, are needed before deployment readiness can be claimed confidently.

Table 85: MRQ5 recoded audit table for Huynh et al. (2025). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
HuynhEtAl2025	2	3	2	3	The review is a systematic review of multi-objective reinforcement learning in smart buildings. Validation-maturity treatment receives score 2 because the paper discusses open-source infrastructures, simulation-oriented environments, and more realistic testbeds, but it does not organize the review as a full validation-maturity taxonomy. Deployment-readiness treatment receives score 3 because the paper repeatedly identifies practical deployment barriers, including computational complexity, long training times, high-quality data needs, heterogeneity across building types and HVAC configurations, interpretability, transparency, and stakeholder trust. Simulation-to-real discussion receives score 2 because the paper discusses simulation/testbed infrastructure and generalization across heterogeneous smart-building settings, but direct field deployment evidence remains weak. Gap-signal visibility receives score 3 because scalability, data requirements, heterogeneity, computational feasibility, interpretability, transparency, and trust are explicit barriers and future directions for real-world MORL adoption.

Table 86: MRQ5 validation-modality scores for Huynh et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Huynh et al. (2025)	3	2	0	0	2

Table 87: MRQ5 validation-rigor scores for Huynh et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Huynh et al. (2025)	2	3	2	2	0

Table 88: MRQ5 gap-signal scores for Huynh et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Huynh et al. (2025)	2	1	2	1	1

Table 89: MRQ5 summary table for Huynh et al. (2025), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Huynh et al. (2025)	Deployment- and generalization-aware, but field-validation thin	Validation is framed mainly through the practical conditions needed for real-world MORL deployment, especially computational feasibility, data sufficiency, robustness, heterogeneity across building types and HVAC configurations, interpretability, transparency, stakeholder trust, and simulation/testbed infrastructure.	The survey contributes to MRQ5 by showing that maturity in MORL for smart buildings depends heavily on whether methods can scale, generalize across diverse buildings, remain understandable and computationally feasible, and move beyond simulation-oriented infrastructure toward stronger real-world validation.

Table 90: One-line MRQ5 answer for Huynh et al. (2025), after recoding.

survey_id	MRQ5 answer
HuynhEtAl2025	Huynh et al. (2025) treat validation maturity mainly as a practical deployment and generalization problem, arguing that MORL for smart buildings still faces barriers in computational complexity, training burden, data needs, cross-building heterogeneity, interpretability, transparency, and trust before it can be considered ready for broader real-world use.

Table 91: MRQ5 recoded audit table for Kannari et al. (2025). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
KannariEtAl2025	3	3	3	3	The review is explicitly focused on reinforcement-learning-based control and optimization in real buildings. Validation-maturity treatment receives score 3 because the paper summarizes practical deployments, implementation settings, deployment modes, and how thoroughly existing applications have been tested. Deployment-readiness treatment receives score 3 because the paper discusses pretraining, static versus dynamic deployment, online relearning, recalibration, safety checks, fallback to traditional control, and validation under varying real-world conditions. Simulation-to-real discussion receives score 3 because the paper explains that real-world reinforcement-learning controllers commonly rely on virtual pretraining before deployment, and it discusses digital twins as data-rich platforms for pre-deployment training and operational adaptation. Gap-signal visibility receives score 3 because the paper explicitly identifies limitations around long-term real-world evidence, scalable deployment, real-world comparability under varying weather and occupancy, and the need for robust validation procedures.

Table 92: MRQ5 validation-modality scores for Kannari et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Kannari et al. (2025)	2	3	2	3	2

Table 93: MRQ5 validation-rigor scores for Kannari et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Kannari et al. (2025)	3	3	3	3	2

Table 94: MRQ5 gap-signal scores for Kannari et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Kannari et al. (2025)	1	0	2	2	3

Table 95: MRQ5 summary table for Kannari et al. (2025), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Kannari et al. (2025)	Strongly deployment- and validation-procedure oriented	Validation is framed through actual real-building implementation practice, including virtual pretraining, static versus dynamic deployment, online relearning, recalibration, safety checks, fallback control, digital twins, and validation under varying real-world conditions.	The survey contributes to MRQ5 by moving beyond general deployment concerns and showing concretely how reinforcement-learning controllers are pretrained, safeguarded, deployed, and validated in real buildings, while also showing that short and heterogeneous test durations leave major uncertainty about long-term reliability, scalability, and comparability.

Table 96: One-line MRQ5 answer for Kannari et al. (2025), after recoding.

survey__id	MRQ5 answer
KannariEtAl2025	Kannari et al. (2025) treat validation maturity mainly as a real-building implementation problem, arguing that reinforcement-learning controllers become more deployment-ready only when they are pretrained safely, equipped with safeguards and fallback logic, validated under varying real-world conditions, and supported by longer-term evidence and digital-twin infrastructure.

17 Michailidis et al.

Table 97: MRQ5 recoded audit table for Michailidis et al. (2025). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
MichailidisEtAl2025	2	2	2	1	The review analyzes reinforcement-learning applications in renewable-energy-integrated building energy management systems using categories that include training methodologies, testbed characteristics, application type, and implementation setting. Validation-maturity treatment receives score 2 because simulated and real-world implementations are made visible through the review design and summary tables, but validation maturity does not structure the paper as a dedicated framework. Deployment-readiness treatment receives score 2 because practical implementation setting is included and both simulated and real-life applications are acknowledged, but the paper does not develop a sustained deployment-readiness critique. Simulation-to-real discussion receives score 2 because the distinction between simulation and real-life application appears in the review categories, although without a strong simulation-to-real methodology discussion. Gap-signal visibility receives score 1 because validation-related gaps are only weakly signaled compared with the paper’s stronger focus on RES integration, reward design, control methods, and application trends.

Table 98: MRQ5 validation-modality scores for Michailidis et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization	evidence
Michailidis et al. (2025)	2	1	1	1	1	

Table 99: MRQ5 validation-rigor scores for Michailidis et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Michailidis et al. (2025)	2	2	2	0	1

Table 100: MRQ5 gap-signal scores for Michailidis et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Michailidis et al. (2025)	1	0	0	0	0

Table 101: MRQ5 summary table for Michailidis et al. (2025), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Michailidis et al. (2025)	Broadly implementation-aware, but validation-light	Validation is framed mainly through application-type, implementation-setting, and testbed descriptors in the review tables, with both simulated and real-life implementations acknowledged, rather than through a strong staged ladder from simulation to field maturity.	The survey contributes to MRQ5 by making implementation setting visible across the reviewed renewable-energy-integrated BEMS literature, but it is weaker than the strongest MRQ5 papers on explicit deployment-readiness critique, safety, long-term validation, and simulation-to-real diagnosis.

Table 102: One-line MRQ5 answer for Michailidis et al. (2025), after recoding.

survey_id	MRQ5 answer
MichailidisEtAl2025	Michailidis et al. (2025) treat validation maturity mainly as an implementation-setting issue by distinguishing simulated and real-life reinforcement-learning applications and including testbed characteristics in their review design, but they contribute more to broad mapping of renewable-energy-integrated BEMS contexts than to a strong framework for deployment readiness or the simulation-to-real gap.

18 Li et al.

Table 103: MRQ5 recoded audit table for Li et al. (2025). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity	treatment	Deployment-readiness	treatment	Simulation-to-real	discussion	Gap-signal	visibility	Recoded evidence	justification
LiEtAl2025	1		3		2		2		Validation-maturity treatment receives score 1 because the survey never classifies the reviewed studies by validation modality and never counts how many were simulated, emulated or deployed. The summary table carries a deployment-context column, but its entries describe the system setting, such as grid-connected residential photovoltaic-battery systems or grid-interactive low energy office buildings, rather than the kind of evidence produced, and the terms testbed, pilot and field test do not occur anywhere in the survey. Deployment-readiness treatment receives score 3 because a dedicated subsection on real-world deployment organizes the barriers into stability and robustness, hardware and computational limits on edge devices, data quality requirements, policy constraints, and integration with proprietary building management platforms where protocol incompatibility requires intermediary control layers with safety fallbacks. Simulation-to-real discussion receives score 2 because the gap is explained at paragraph level in three separate places, through the statement that discrepancies between simulated and real environments cause models to overfit and lose performance, through sim-to-real transfer risk named as an open challenge in safe reinforcement learning, and through digital twins proposed as a pre-training route to reduce transfer gaps. Gap-signal visibility receives score 2 because the survey states that future research must address the identified deployment constraints and offers mitigation strategies including multi-scenario training, Bayesian reinforcement learning and confidence intervals on policy evaluation, without developing a validation research agenda.	

Table 104: MRQ5 validation-modality scores for Li et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization	evidence
Li et al. (2025)	2	0	0	1	1	

Table 105: MRQ5 validation-rigor scores for Li et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy	clarity	Deployment-readiness	discussion	Simulation-to-real gap	rigor	Safety / safeguards for deployment	Strength of real-world evidence
Li et al. (2025)	1		3		2		3	1

Table 106: MRQ5 gap-signal scores for Li et al. (2025), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Li et al. (2025)	1	0	1	2	0

Table 107: MRQ5 summary table for Li et al. (2025), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Li et al. (2025)	Simulation-oriented, deployment-aware	The survey frames validation as an engineering obstacle rather than as an evidence question. Simulation is treated as the environment in which learning happens and whose fidelity must be improved, while deployment is treated as a separate downstream problem of hardware limits, protocol integration and safety fallbacks. The reviewed studies are never sorted by the kind of evidence they produce.	For MRQ5, the paper shows a survey that discusses deployment barriers in concrete detail while leaving the underlying evidence base uncharacterized. Because no study is classified as simulated, emulated or deployed, the survey cannot say how much of the reported performance rests on simulation, and it correspondingly identifies no shortage of pilots or field tests and no shortage of long-term real-world evidence. Safety receives substantive treatment through safe reinforcement learning, which is the strongest deployment-facing element of the review.

Table 108: One-line MRQ5 answer for Li et al. (2025), after recoding.

survey_id	MRQ5 answer
LiEtAl2025	Li et al. (2025) devote a dedicated subsection to real-world deployment barriers and give safety a substantive treatment through safe reinforcement learning, but never classify the reviewed studies by the kind of validation evidence they produce; simulation fidelity and transfer risk are discussed as engineering problems to be mitigated rather than as limits on what the literature has demonstrated, and the survey signals no shortage of pilots, field tests or long-term evidence.

Table 109: MRQ5 recoded audit table for Togashi (2025). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
Togashi2025	1	1	2	2	Validation-maturity treatment receives score 1 because the survey never classifies the seventy-nine reviewed studies by the kind of evidence they produce, and its exclusion criteria actively remove work on reinforcement learning testbeds from the corpus on the grounds that such studies concern testbed development rather than reinforcement learning application. Deployment-readiness treatment receives score 1 because deployment is never discussed as a topic: the term does not occur in the survey, and real buildings are referenced once, only as the contrast case against simulation. Simulation-to-real discussion receives score 2 because the gap is explained concretely in two places. The survey argues that few reward functions include a control-stabilization term because simulation-based studies do not reproduce the equipment degradation caused by frequent on-off cycling that occurs in real buildings, so the problem never becomes apparent, and because the simulations used are often not at minute or second resolution, with discrete static simulations still common. Separately it argues that reward functions embedding solution hints allow agents to perform well in simulation while failing to generalize. Gap-signal visibility receives score 2 because the survey concludes that control stabilization will need to be incorporated into reward functions in future work and warns that the field risks generating case studies that are difficult to compare, but it develops no validation or deployment research agenda.

Table 110: MRQ5 validation-modality scores for Togashi (2025), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Togashi (2025)	2	1	0	1	1

Table 111: MRQ5 validation-rigor scores for Togashi (2025), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Togashi (2025)	0	1	2	0	1

Table 112: MRQ5 gap-signal scores for Togashi (2025), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Togashi (2025)	2	0	1	0	1

Table 113: MRQ5 summary table for Togashi (2025), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Togashi (2025)	Simulation-implicit, validation not characterized	Simulation is treated as the unexamined default setting in which reinforcement learning research takes place. The survey turns to it only when simulation practice distorts what it is analysing, arguing that the absence of equipment degradation and the coarse temporal resolution of common simulations explain why control stabilization is largely missing from reward functions. Validation modality is never a property the reviewed studies are sorted by.	For MRQ5, the paper contributes an unusually specific diagnosis without an accompanying agenda. It identifies exactly how simulation practice suppresses a design concern, naming both the missing physical mechanism and the missing temporal resolution, yet it never asks how much of the literature is simulation-based, never discusses deployment, and does not raise safety at all. The scope decision to exclude testbed-development studies removes from the corpus the very work that would characterize the evidence base.

Table 114: One-line MRQ5 answer for Togashi (2025), after recoding.

survey_id	MRQ5 answer
Togashi2025	Togashi (2025) treats simulation as the unexamined default and never classifies the seventy-nine reviewed studies by validation modality, having excluded testbed-development work from the corpus by design, yet it supplies one of the sharpest simulation-to-real observations in the corpus: control stabilization is largely absent from reward functions because simulations do not reproduce equipment degradation from frequent cycling and are often too coarse in time resolution for cycling to be visible at all.

Table 115: MRQ5 recoded audit table for Lei et al. (2026). Scores use the shared 0–3 evidence-intensity scale.

survey_id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
LeiEtAl2026	3	3	3	3	Validation-maturity treatment receives score 3 because the survey counts its own evidence base rather than describing it: of the fourteen studies in the application table, thirteen compare against a conventional baseline and all of those report superior or equivalent performance on at least one resilience-relevant metric, five provide field-validated evidence from real occupied buildings, the remainder rely on simulation environments often calibrated with real-world datasets, and eight address a resilience-relevant disturbance scenario such as grid outage, supply interruption, extreme weather or climate variability. Reported reductions in energy cost or consumption range from approximately 9.7 to 45 percent relative to the best-performing conventional baseline, and reductions in thermal comfort violations reach 43 percent in real-building installations. Deployment-readiness treatment receives score 3 because a dedicated subsection is given to emerging real-world application cases with named installations and quantified outcomes. A soft actor–critic controller was deployed in the thermally activated building system of a real office building at EMPA NEST in Switzerland, pre-trained on a calibrated resistance-capacitance thermal model and transferred directly to the physical system, where it reduced temperature violations by 68 percent relative to the best-performing rule-based baseline without increasing energy consumption. At district scale, a model-free reinforcement learning framework was validated across a real multi-building district heating network, demonstrating cross-building demand flexibility without explicit physical models of individual buildings. Deployment architecture is developed further through decentralized fallback control and delay-aware policy design. Simulation-to-real discussion receives score 3 because transfer is treated across several parts of the survey: one deployment is described as pre-trained on a calibrated thermal model and transferred directly to the physical system, a dedicated challenge subsection explains that policies trained in one building may degrade or become unsafe when transferred to another and that agents overfit implicit occupancy patterns, and simulation environments are said to be unable to faithfully reproduce the full complexity of cascading failures. Gap-signal visibility receives score 3 because the survey states that field validations specifically targeting resilience under high-impact low-probability events remain scarce and represent a critical frontier for future empirical investigation, and separately identifies heterogeneity in system scale, energy carrier composition, baseline choice and metric definition as grounds for standardized benchmarking protocols.

Table 116: MRQ5 validation-modality scores for Lei et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Lei et al. (2026)	3	2	2	3	3

Table 117: MRQ5 validation-rigor scores for Lei et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Lei et al. (2026)	2	3	2	3	2

Table 118: MRQ5 gap-signal scores for Lei et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Lei et al. (2026)	3	3	2	2	2

Table 119: MRQ5 summary table for Lei et al. (2026), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Lei et al. (2026)	Evidence-counting, deployment-forward	The survey frames validation as a property of the literature that can be measured. Studies are sorted by whether they were field-validated in occupied buildings or run in simulation, whether they compared against a conventional baseline, and whether they addressed a resilience-relevant disturbance, and the counts are reported. Deployment is treated as an active frontier with named installations rather than as a distant prospect.	For MRQ5, this is the corpus case where a survey quantifies the balance between simulated and real evidence and then draws the limiting conclusion from it. Reporting five field-validated studies out of fourteen, with energy reductions spanning roughly 9.7 to 45 percent and comfort violation reductions reaching 43 percent in real buildings, lets the survey argue specifically that deployment feasibility is established while resilience under high-impact low-probability conditions is not, and that the scarcity of such field validation is the critical frontier. Safety is developed alongside this through hard-constraint methods, barrier certificates and model-predictive filtering, which ties deployment readiness to verifiable guarantees.

Table 120: One-line MRQ5 answer for Lei et al. (2026), after recoding.

survey__id	MRQ5 answer
LeiEtAl2026	Lei et al. (2026) count the evidence base rather than characterizing it in general terms, reporting that five of fourteen reviewed applications are field-validated in real occupied buildings and thirteen compare against a conventional baseline, and naming two deployments with quantified outcomes; on that basis the survey argues that transfer from simulation to real infrastructure is technically demonstrated while field validation under high-impact low-probability conditions remains scarce, and pairs this with substantive treatment of safety guarantees for deployment.

21 Mazhar et al.

Table 121: MRQ5 recoded audit table for Mazhar et al. (2026). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
MazharEtAl2026	2	3	1	2	Validation-maturity treatment receives score 2 because the survey discusses the validation landscape at paragraph level and names concrete installations, including a net-zero office building in Amsterdam using real-time sensing with deep reinforcement learning to optimize heating, ventilation, air conditioning, lighting and renewable distribution, and an intelligent campus in France regulating indoor climate from occupancy data, but it never classifies or counts the reviewed studies by the kind of evidence they produce, and the named cases are reported without quantitative outcomes. Deployment-readiness treatment receives score 3 because a dedicated subsection sets out a three-phase roadmap from pilot testing in controlled single-zone settings, through incremental scaling to multi-zone buildings addressing computational overhead and interoperability, to infrastructure optimization, with key challenges and resource-constrained-setting considerations attached to each phase, and because a further subsection addresses the economic barriers to adoption. Simulation-to-real discussion receives score 1 because the transfer problem is never raised: the term sim-to-real does not occur, simulation is mentioned almost exclusively to name platforms rather than to discuss fidelity, and the survey nowhere considers how results obtained in simulation carry over to operating buildings. Gap-signal visibility receives score 2 because the lack of real-world benchmarks is named among the stated research gaps and the datasets subsection closes by calling for large-scale real-time data from multi-zone buildings, but the gaps identified concern infrastructure availability rather than the strength of the existing evidence.

Table 122: MRQ5 validation-modality scores for Mazhar et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Mazhar et al. (2026)	2	2	2	2	1

Table 123: MRQ5 validation-rigor scores for Mazhar et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Mazhar et al. (2026)	1	3	1	1	2

Table 124: MRQ5 gap-signal scores for Mazhar et al. (2026), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Mazhar et al. (2026)	1	2	1	2	1

Table 125: MRQ5 summary table for Mazhar et al. (2026), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Mazhar et al. (2026)	Deployment-oriented, evidence uncounted	The survey frames validation as a question of available infrastructure and adoption pathway rather than of evidential strength. Datasets and benchmarking environments are catalogued as resources for testing, and the route to deployment is laid out as a staged programme running from pilot testing through incremental scaling to infrastructure optimization, with attention to cost and to adoption in resource-constrained settings.	For MRQ5, the paper is the corpus case where deployment is planned in detail while the evidence base behind it stays uncharacterized. Named installations are offered as proof of practical significance without reported figures, no study is classified as simulated or deployed, and the transfer problem is absent entirely, so the roadmap rests on an assumption of transferability that the survey never examines.

Table 126: One-line MRQ5 answer for Mazhar et al. (2026), after recoding.

survey__id	MRQ5 answer
MazharEtAl2026	Mazhar et al. (2026) give deployment the most detailed treatment of any element of the review, setting out a three-phase roadmap from pilot testing to infrastructure optimization and cataloguing the datasets and benchmarking environments available for validation, but the evidence base itself is never characterized: named installations are cited without outcomes, no study is classified by validation modality, and the simulation-to-real transfer problem is not raised anywhere in the survey.

22 Mohsenpour and Xing

Table 127: MRQ5 recoded audit table for Mohsenpour and Xing (2026). Scores use the shared 0–3 evidence-intensity scale.

survey__id	Validation-maturity treatment	Deployment-readiness treatment	Simulation-to-real discussion	Gap-signal visibility	Recoded evidence justification
MohsenpourXing2026	3	3	3	3	Validation-maturity treatment receives score 3 because a dedicated subsection separates the reviewed work by the kind of evidence it produces, opening with the finding that relatively few studies have progressed beyond simulation-based evaluation, that most validate controllers exclusively within simulation environments where results do not capture the complexities and uncertainties of operational buildings, and that most real-world studies prioritize energy efficiency while overlooking indoor environmental quality and thermal comfort outcomes. Deployment-readiness treatment receives score 3 because four field deployments are reported with quantified outcomes, including cost reductions of 14 and 16 percent across two thermal zones under multi-agent control, a 13.9 percent reduction in cooling energy with an 11 percent improvement in thermal acceptability in a 164 square metre net-zero office, and a 68 percent improvement in indoor temperature regulation over a two-month cooling season without increased energy consumption, and because an implementation strategy figure accompanies the analysis. Simulation-to-real discussion receives score 3 because the transfer route is described mechanically rather than asserted: one deployment pre-trains offline on a building simulation library, then performs continuous online learning on daily thermal comfort, indoor conditions and occupant complaints, then transitions to full deployment after roughly five to ten weeks without further virtual training data, while another trains a surrogate model on simulation output because full simulation is computationally expensive, and real-world performance is reported against simulation-based performance. Gap-signal visibility receives score 3 because the survey states the need for studies extending reinforcement learning to real-world settings and separately identifies that existing real-world work neglects indoor environmental quality, which makes the gap specific rather than general.

Table 128: MRQ5 validation-modality scores for Mohsenpour and Xing (2026), recoded using the shared evidence-intensity scale.

Survey	Simulation	Emulation / testbed	Pilot / demonstration	Real building / field deployment	Cross-site / transfer / generalization evidence
Mohsenpour and Xing (2026)	3	2	2	3	3

Table 129: MRQ5 validation-rigor scores for Mohsenpour and Xing (2026), recoded using the shared evidence-intensity scale.

Survey	Validation taxonomy clarity	Deployment-readiness discussion	Simulation-to-real gap rigor	Safety / safeguards for deployment	Strength of real-world evidence
Mohsenpour and Xing (2026)	2	3	3	1	3

Table 130: MRQ5 gap-signal scores for Mohsenpour and Xing (2026), recoded using the shared evidence-intensity scale.

Survey	Overreliance on simulation	Lack of pilots / field tests	Weak sim-to-real methodology	Limited deployment protocol detail	Limited long-term real-world evidence
Mohsenpour and Xing (2026)	3	3	2	2	2

Table 131: MRQ5 summary table for Mohsenpour and Xing (2026), after recoding.

Survey	Validation orientation	Dominant validation framing	Main MRQ5 consequence
Mohsenpour and Xing (2026)	Deployment-evidenced, transfer-explicit	The survey frames validation as a threshold the literature has mostly not crossed, and treats the crossing itself as the object of analysis. Rather than contrasting simulation with reality in the abstract, it reports how specific controllers made the transition, through offline pre-training on building simulation libraries, surrogate models substituting for expensive simulation, and staged online learning that ingests occupant complaints before the system operates independently.	For MRQ5, this is the corpus case that pairs a quantified deployment record with a second-order critique of it. The survey reports field results in real buildings with figures, then argues that the small body of real-world work is itself skewed, because those studies optimize energy while neglecting the indoor environmental quality and comfort outcomes the review treats as co-equal. The evidence gap is therefore not only that too few studies leave simulation, but that the ones which do measure the wrong things. Safety and safeguards receive little attention by comparison.

Table 132: One-line MRQ5 answer for Mohsenpour and Xing (2026), after recoding.

survey_id	MRQ5 answer
MohsenpourXing2026	Mohsenpour and Xing (2026) devote a dedicated subsection to real-world applications, report four field deployments with quantified outcomes, and describe the transfer route mechanically through offline pre-training, surrogate models and staged online learning that ingests occupant complaints before full autonomy; the survey then adds a second-order finding, that the few studies leaving simulation optimize energy while neglecting indoor environmental quality and comfort, making the evidence gap one of measurement as well as of quantity.